

Learning Equilibria in Stochastic Information Flow Tracking Games with Partial Knowledge

Shruti Misra^{*}, Shana Moothedath^{*}, Hossein Hosseini^{*}, Joey Allen[†], Linda
Bushnell^{*}, Wenke Lee[†], and Radha Poovendran^{*}

^{*}Network Security Lab, Department of Electrical and Computer
Engineering, University of Washington, Seattle, WA, USA

[†]School of Computer Science, College of Computing, the Georgia
Institute of Technology, Atlanta, GA, USA

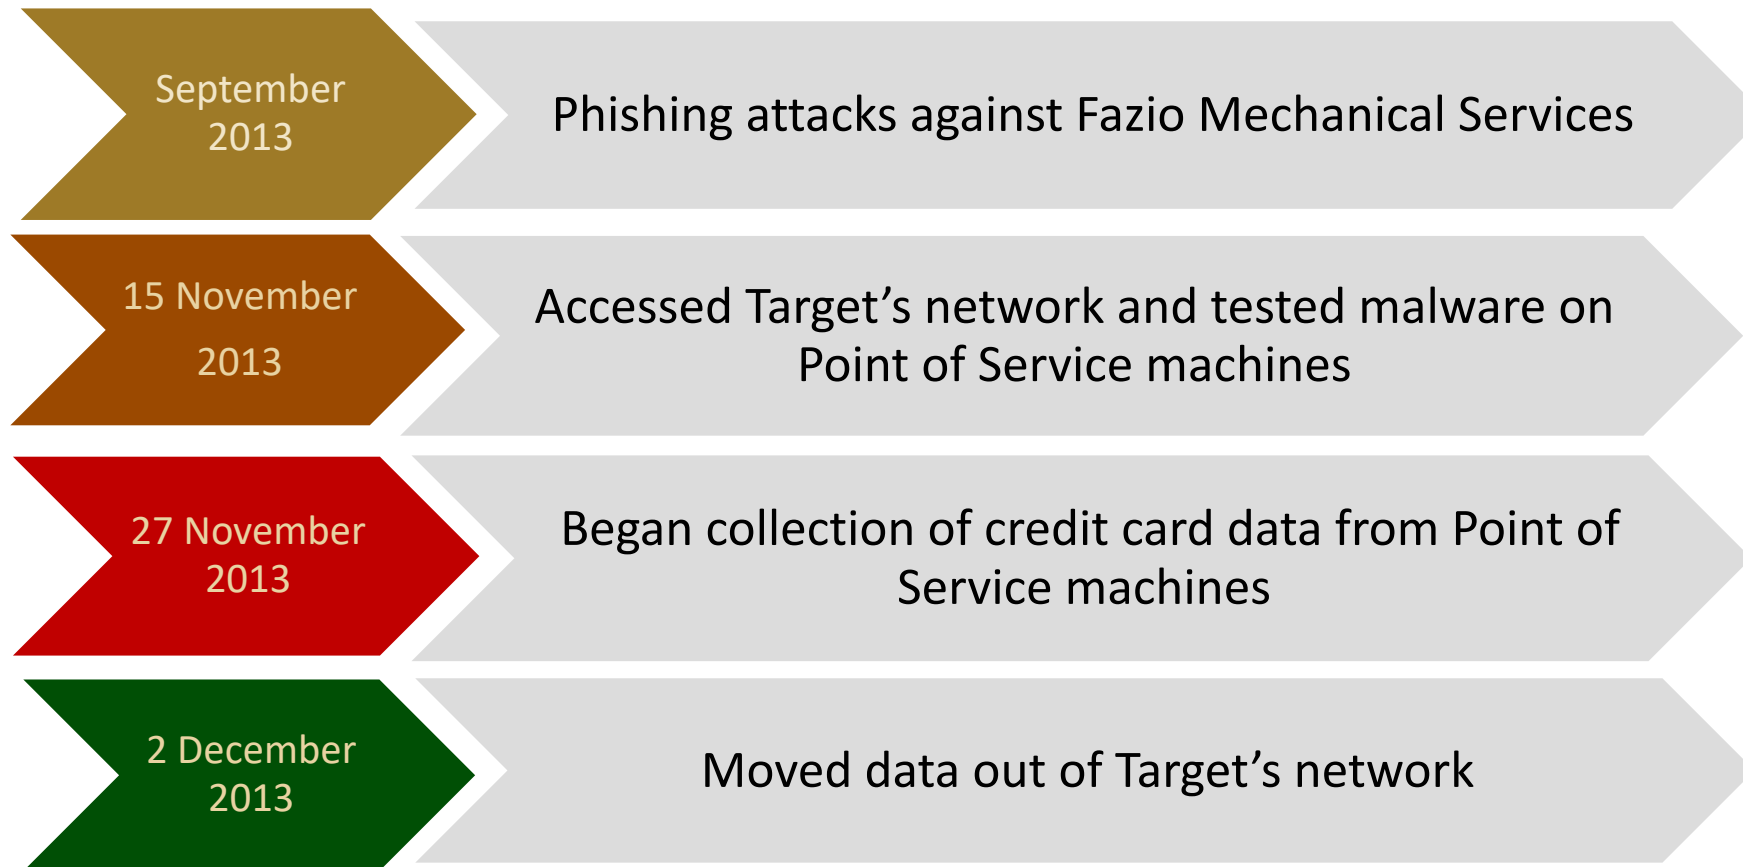
Advanced Persistent Threats

An advanced persistent threat (APT) is a prolonged and targeted cyberattack in which an intruder gains access to a network and remains undetected for a period of time.

The intention of an APT attack is usually to monitor network activity and steal data rather than to cause damage to the network or organization.

2013 Target Corporation Data Breach

40 million credit and debit card numbers and 70 million records of personal information were stolen

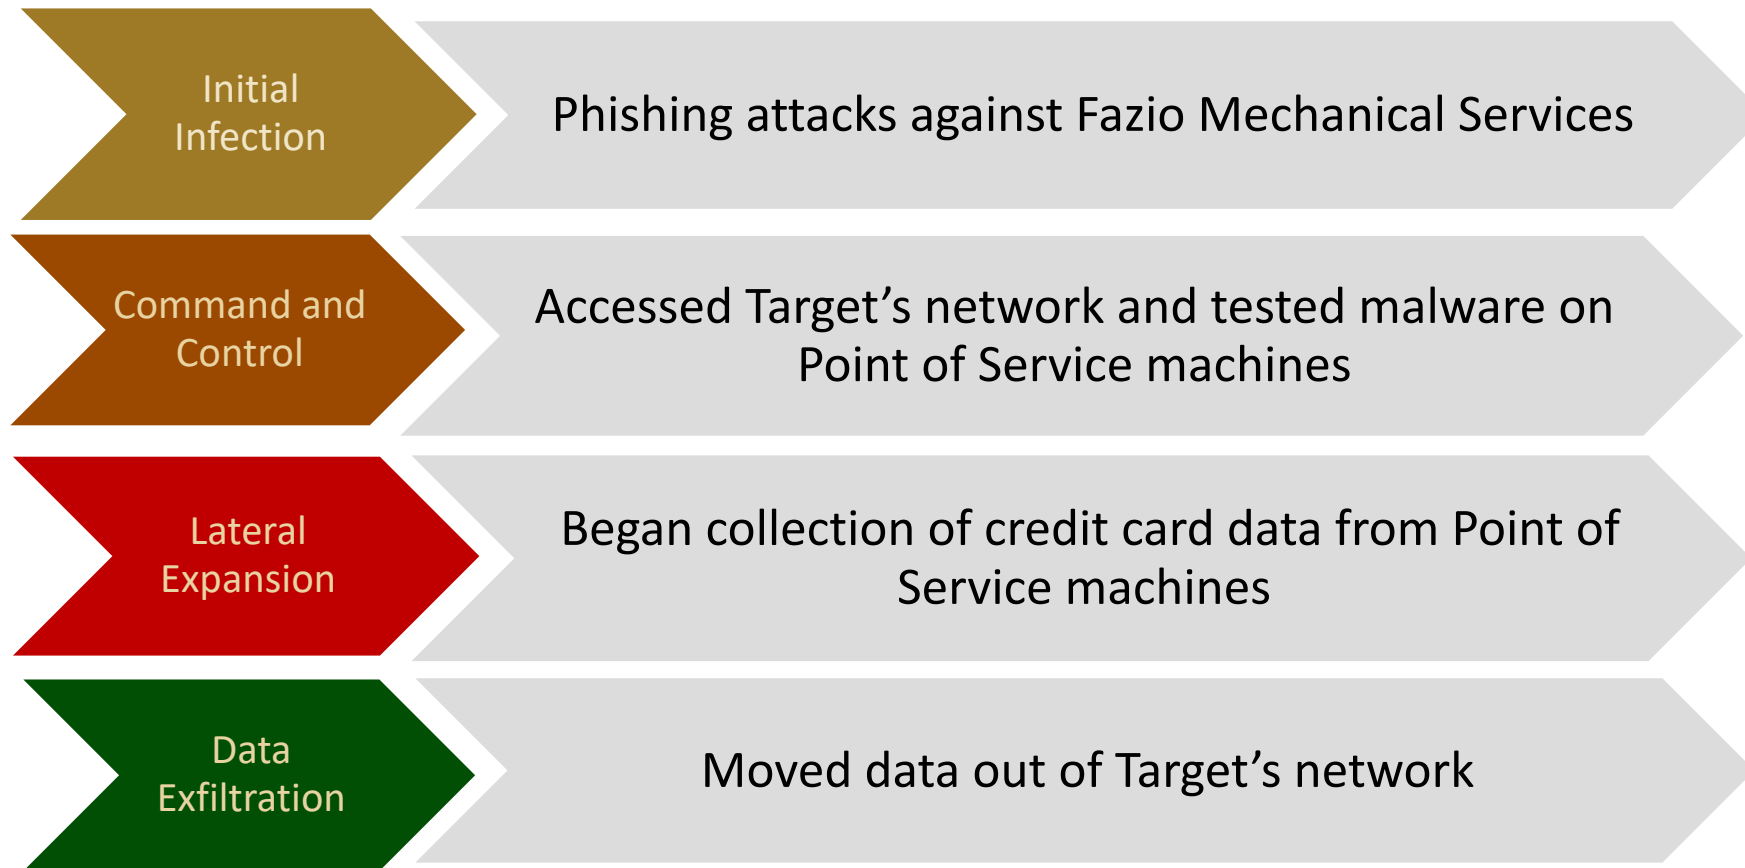


Shu, X., Tian, K., Ciambrone, A. and Yao, D., 2017. Breaking the target: An analysis of target data breach and lessons learned. *arXiv preprint arXiv:1701.04940*.



Stages of Advanced Persistent Threats (APTs)

40 million credit and debit card numbers and 70 million records of personal information were stolen



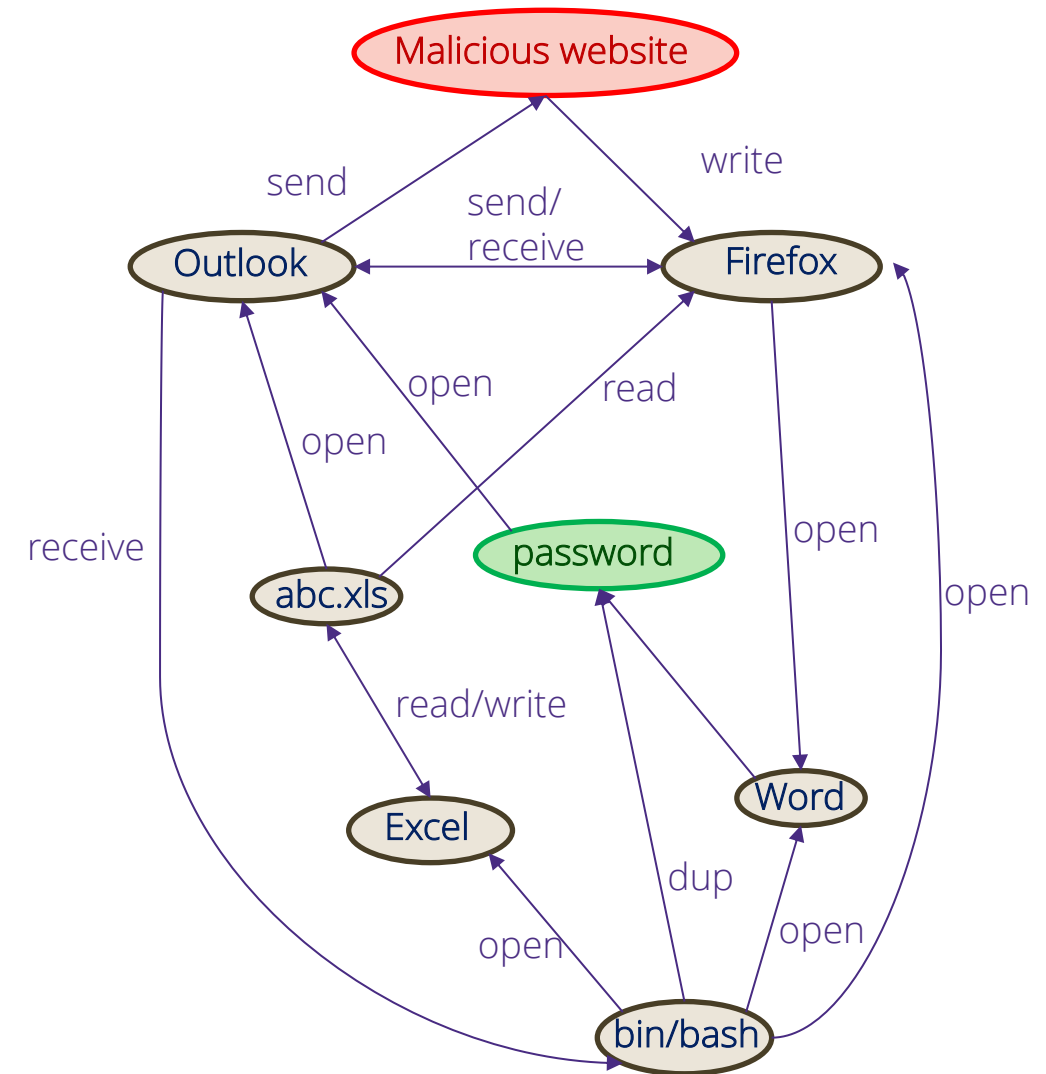
Shu, X., Tian, K., Ciabrone, A. and Yao, D., 2017. Breaking the target: An analysis of target data breach and lessons learned. *arXiv preprint arXiv:1701.04940*.



Information Flow Graphs (IFGs)

Adversary interact with system components (e.g. processes, files) using **system calls** (e.g. read, write, open)

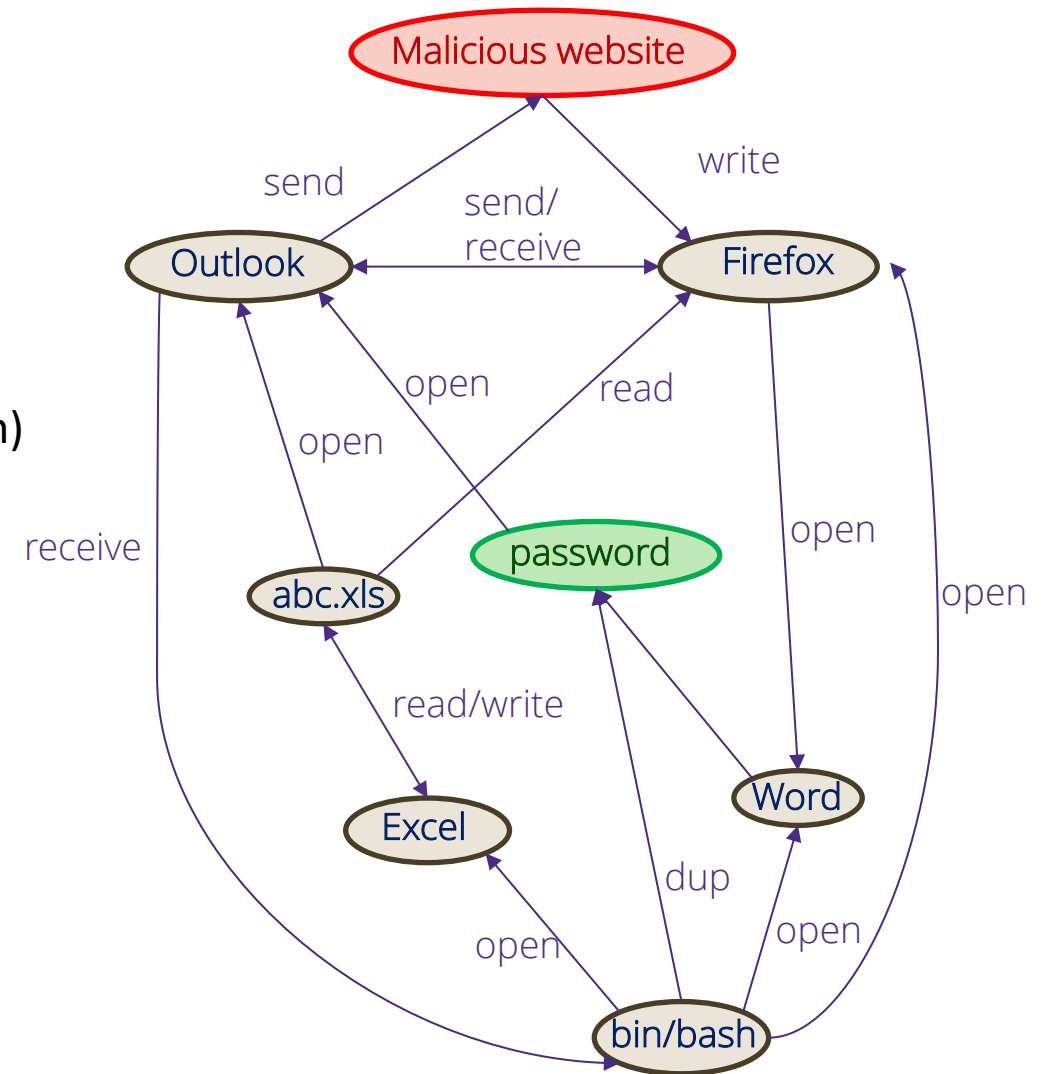
Information flows (data- and control-flow commands) abstract the interaction of the adversary with various system components.



Information Flow Graphs (IFGs)

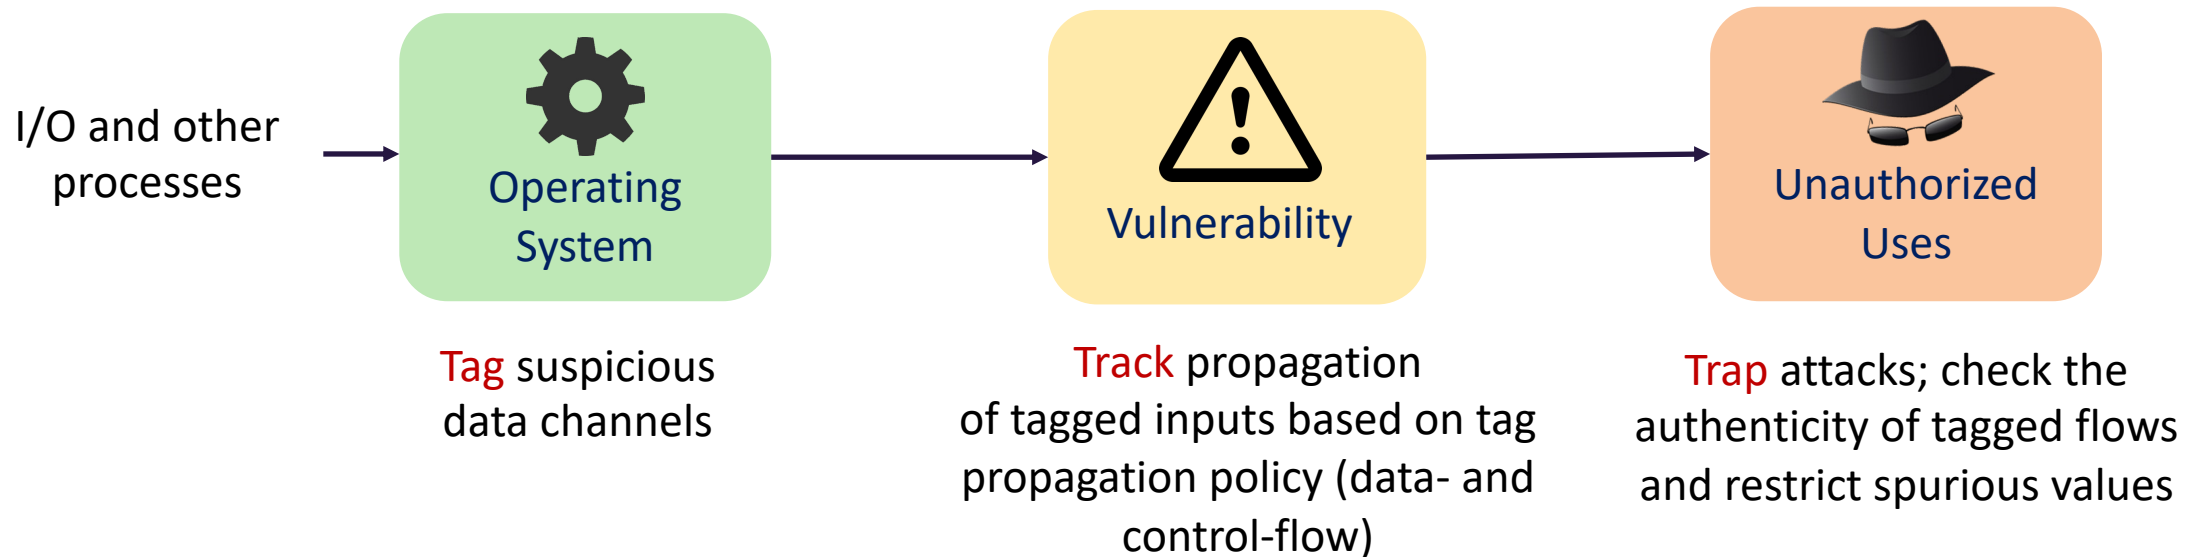
Graphical representation of system log data

- **Nodes:** Subjects and Objects
 - Subjects: Processes (an instance of a computer program)
 - Objects: e.g. files/memory/network sockets
- **Edges:** probability of transferring information flows between processes
 - System calls: e.g. read, write, open, send



Dynamic Information Flow Tracking (DIFT)

DIFT is a flow tracking-based mechanism that is widely used to detect APTs.



Suh, G. E., Lee, J. W., Zhang, D., & Devadas, S. (2004, October). Secure program execution via dynamic information flow tracking. *In ACM Sigplan Notices*, ACM.

IEEE Control and Decision Conference 2019

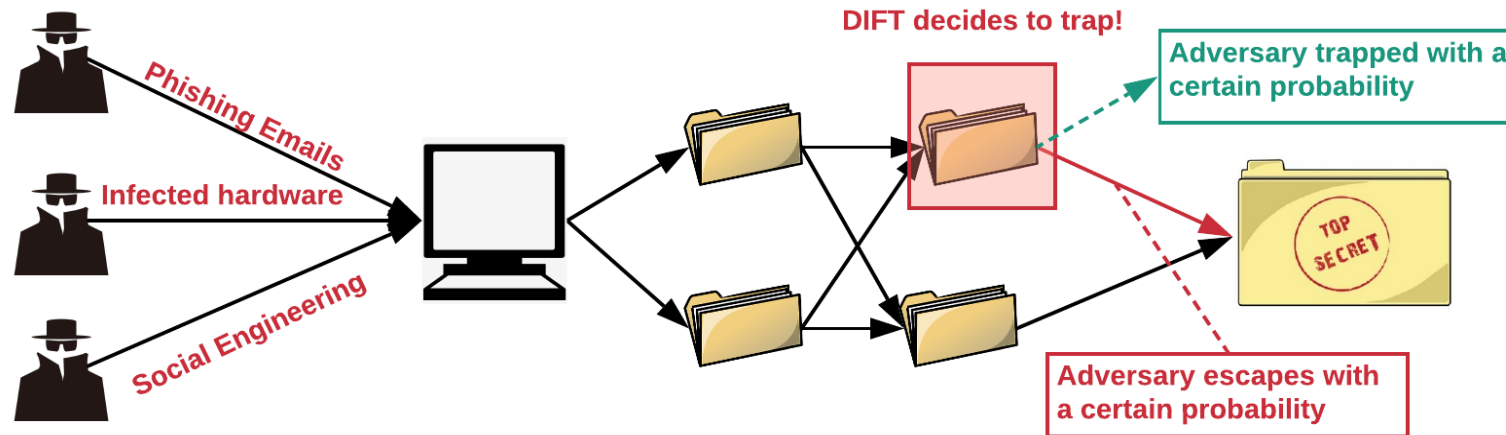


Dynamic Information Flow Tracking (DIFT)

Implementation and operation of DIFT, however, introduce memory and performance overhead on the system as it involves tracking and analyzing a large number of benign flows [3]. Thus an optimal selection of processes in the system to perform security analysis is critical for effective and resource efficient detection.



Problem Formulation



Aim


Model a cost-effective DIFT-based defense mechanism against APTs that:

- a. Captures the trade-off between detection accuracy and resource efficiency.
- b. Accounts for rate of false negatives.


Problem Formulation

- APT (**A**) is the adversary infecting the system.
- DIFT (**D**) is the defense framework used against **A**.
- Goal of **A** is to **evade** detection and reach target nodes.
- Goal of **D** is to dynamically place **trap locations** to:
 - Maximize **detection probability**
 - Minimize **performance overhead** on system
- **Strategic interactions** between **D** and **A** form a **game**.

State space and Actions



Information Flow
Graph



State Space

τ : Trap State
 ϕ : Drop State

$$\mathbf{S} = \{s_0, s_1, \dots, s_{N-1}, d, \phi, \tau\}$$

Collection of
entry nodes

Node in
IFGs

Terminal
state

Quit
state

Trapped
state

We assume that both players know destination node d .

ACTIONS:

- $\mathbf{A}(a_t)$: {Transitioning to a out-neighboring state, Quit}
- $\mathbf{D}(d_t)$: {Trap, no Trap}
- No actions are allowed for D at nodes in the set $\{s_0, \phi, \tau\}$

Transition Structure

DIFT (D)

$$p_D = \{p_D(s) : s \in \mathbf{S}, s \notin \{s_0, d, \phi, \tau\}\}$$

$p_D(s)$: Probability of trapping in s

$1 - p_D(s)$: Probability of not trapping in s

False negatives of DIFT ($FN(s_i)$) depends on:

- Number of security rules
- Strength security rules

We assume policies of both players to be stationary

APT (A)

For $s \notin \{s_0, d, \tau, \phi\}$

$$p_A(s) = \{p(s, s') : s \in \mathbf{S}, s' \in N(s)\}$$

$p_A(s)$: Probability distribution of all possible actions in state s .

Given actions $a_t = s_{i'}$ and $d_t = 1$. The next state is given by:

$$s_{t+1} = \begin{cases} s_{i'} & \text{w.p. } FN(s_i) \\ \tau & \text{w.p. } 1 - FN(s_i) \end{cases}$$



Payoff Structure

DIFT (D)

$$r_D(s, d, a) \begin{cases} \alpha_D > 0, \text{ if } s = \tau & \text{Reward for detecting} \\ \beta_D < 0, \text{ if } s = d & \text{Penalty for evading detection} \\ \sigma_D \geq 0, \text{ if } s = \phi & \text{Reward for dropping out} \\ C_D(s) < 0, \text{ if } s \in S \text{ and } d = 1 & \text{Resource cost} \\ 0, \text{ otherwise} \end{cases}$$

APT (A)

$$r_A(s, d, a) \begin{cases} \alpha_A < 0, \text{ if } s = \tau & \text{Penalty for getting detected} \\ \beta_A > 0, \text{ if } s = d & \text{Reward for reaching destination} \\ \sigma_A \leq 0, \text{ if } s = \phi & \text{Penalty for dropping out} \\ 0, \text{ otherwise} \end{cases}$$



Payoff Structure

$$U_A(p_D, p_A) = p_{\tau}^{(T)} \alpha_A + p_R^{(T)}(d) \beta_A + p_{\phi}^{(T)} \sigma_A$$

$$U_D(p_D, p_A) = \sum_{v_i \in V(G)} (p_D(s_i) \mathcal{C}_D(v_i)) + p_{\tau}^{(T)} \alpha_D + p_R^{(T)}(d) \beta_D + p_{\phi}^{(T)} \sigma_D$$

$p_{\tau}^{(T)}$: Probability of being in the trapped state at terminal time T

$p_R^{(T)}$: Probability of reaching the destination at terminal time T

$p_{\phi}^{(T)}$: Probability of being drop state at terminal time T

Terminating Conditions

1. The adversary is trapped by the defender.
2. The adversary reaches the destination.
3. The adversary quits the game.



Payoff Structure

$$U_A(p_D, p_A) = p_\tau^{(T)} \alpha_A + p_R^{(T)}(d) \beta_A + p_\phi^{(T)} \sigma_A$$

$$U_D(p_D, p_A) = \sum_{v_i \in V(G)} (p_D(s_i) \mathcal{C}_D(v_i)) + p_\tau^{(T)} \alpha_D + p_R^{(T)}(d) \beta_D + p_\phi^{(T)} \sigma_D$$

$$p_\tau(s_i) = p_D(s_i) F N(s_i)$$

$$p_R(s_i) = \sum_{s_j \in S} p_A(s_j, s_i) p_R(s_i) (1 - p_\tau(s_j))$$

$$p_\phi(s_i) = p_A(s_i, \phi) p_R(s_i) (1 - p_\tau(s_j))$$

Terminating Conditions

1. The adversary is trapped by the defender.
2. The adversary reaches the destination.
3. The adversary quits the game.



Properties of APT vs. DIFT game

Nonzero-sum

Defender payoff:

Reward, penalty,
cost of tagging

Adversary
payoff:

Reward, penalty

Stochastic

Uncertainty in
both players'
actions

Imperfect Information

Defender cannot
distinguish between
benign and
malicious flow

Incomplete Information

Unknown rate of
false negatives

Solution Concept

Strategy pair $(p_A^*, p_D^*) \in \mathbf{P}_A \times \mathbf{P}_D$ forms a Nash equilibrium in stochastic stationary strategies for any $\epsilon > 0$ and for all $p_A \in \mathbf{p}_A$ and $p_D \in \mathbf{p}_D$ if

$$\begin{aligned} U_A(p_A^*, p_D^*) &\geq (p_A, p_D^*) \\ U_D(p_A^*, p_D^*) &\geq (p_A^*, p_D) \end{aligned}$$

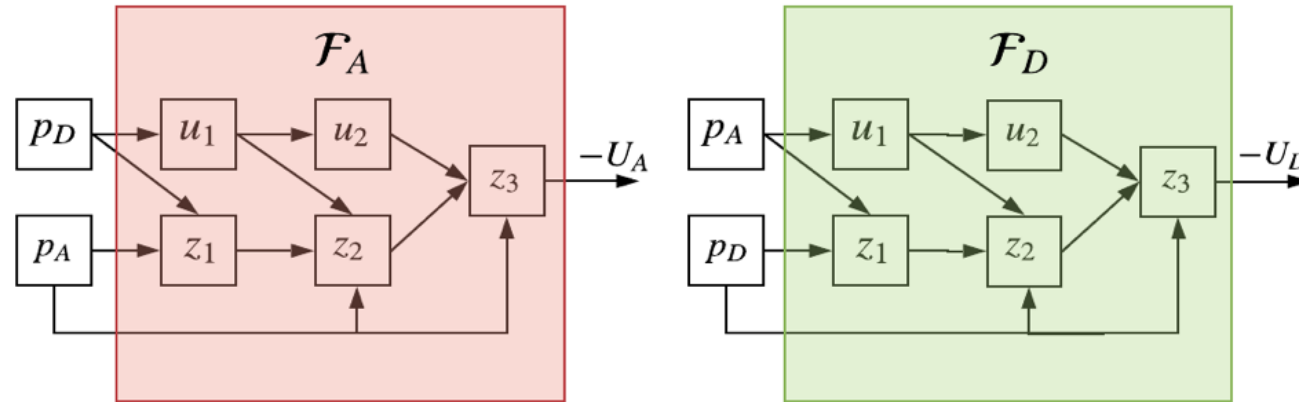
Result 1:

The APT vs. DIFT game satisfies the following:

1. The game terminates in at most $N + 3$ number of steps.
2. There exists a Nash Equilibrium (NE) for the game.
3. For a given strategy pair (p_A, p_D) computation of $p_R^{(T)}$, $p_\tau^{(T)}$ and $p_\phi^{(T)}$ has complexity **linear** in the number of states and edges in S .



Modified Partially Input Convex Neural Networks (PICNNs)



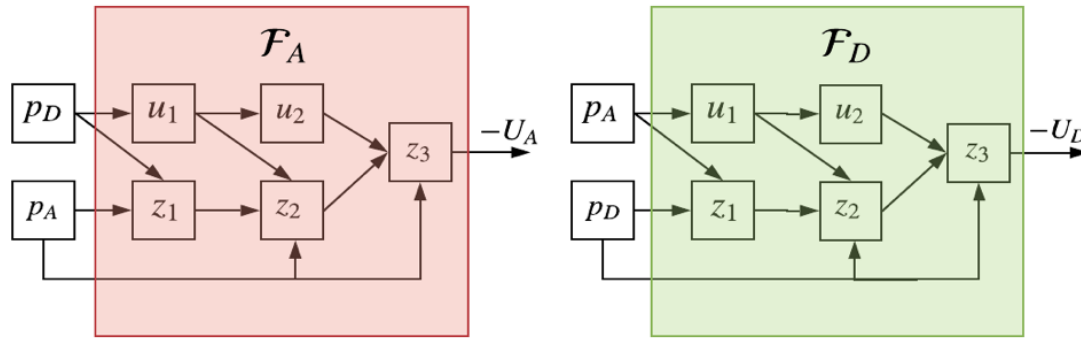
The modified PICNN model defines a neural network with k layers over output y by implementing the following architecture for $i = 0, \dots, k$

$$z_{i+1} = g_i \left(W_i^{(z)} z_i + W_i^{(y)} y + W_i^{(u)} u_i + b_i \right), f(y; \theta) = z_k, u_0 = x$$

B. Amos, L. Xu, and J. Z. Kolter, "Input convex neural networks," in Proceedings of the 34th International Conference on Machine Learning, vol. 70. JMLR. org, 2017, pp. 146–155.



Results



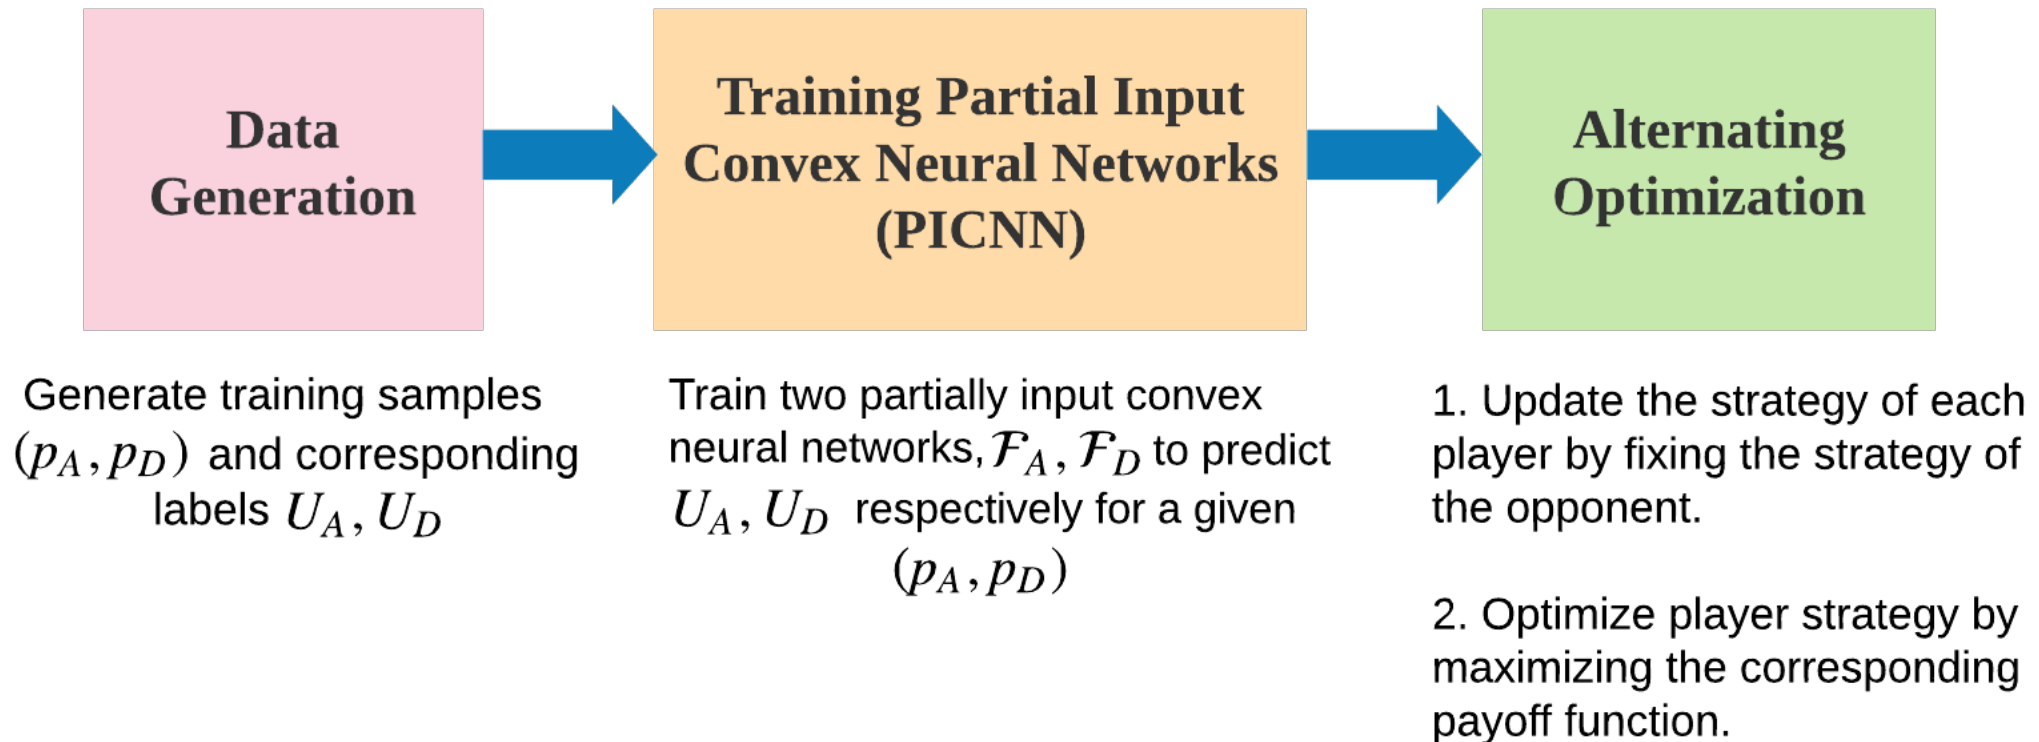
$$z_{i+1} = g_i \left(W_i^{(z)} z_i + W_i^{(y)} y + W_i^{(u)} u_i + b_i \right), f(y; \theta) = z_k, u_0 = x$$

Result 2:

The function \mathcal{F} is convex in y provided that all $W_{(1:k-1)}^{(z)}$ are non-negative, and all functions g_i are convex and non-decreasing.




Supervised Learning for Games



Case Studies

To demonstrate the performance of algorithm, we conducted two different simulation studies: (i) using a randomly generated IFG and (ii) using an IFG of ScreenGrab attack data obtained by the Refinable Attack Investigation System (RAIN)

Case Study: Synthetic Graph



- Number of nodes = 50
- Erdős Renyi graph with directed edge probability $p = 0.2$
- Destination node (d) is the 50th node.
- Resource cost is generated from a uniform distribution $[0,10]$.

Case Studies

We conducted a sensitivity analysis to validate that the converged strategies correspond to a Nash Equilibrium of the game.


We first initialize the strategies of both players to the strategy returned by algorithm. Then we perturb the adversary's strategy while keeping the defender's strategy fixed. Similarly, we perturb the defender's strategy while keeping the adversary strategy fixed. Here, the rationale is that if the strategies returned by algorithm is a NE (with respect to the approximated payoff functions), then perturbing the strategy of a player should not improve its payoff. Therefore, the payoff from perturbed strategies should be equal to or less than that of the output of the algorithm.

The results provided below are obtained by generating 100 different perturbations of each player's strategy while keeping other player's strategy fixed.




Simulation Results

APT's payoff vs. iteration




Payoff Convergence

APT's payoff for 100 perturbed strategies



Convex Approximation


APT's payoff for 100 perturbed strategies



Original Function


Simulation Results

DIFT's payoff vs. iteration



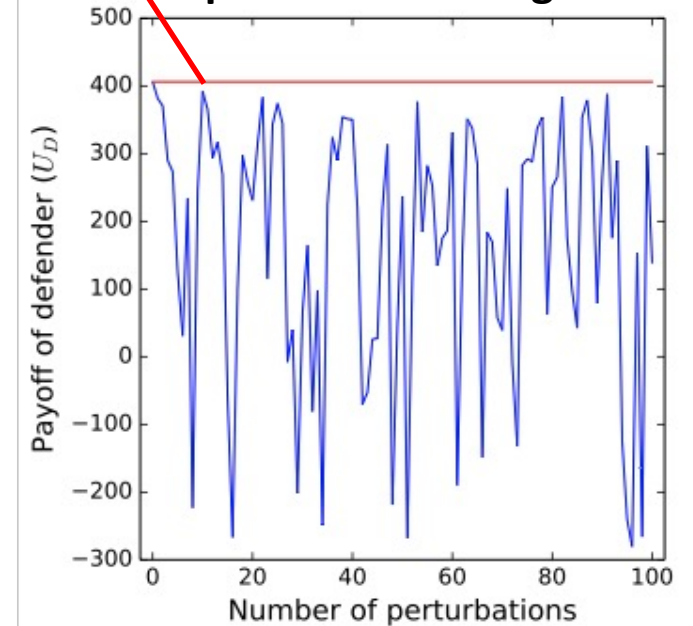
Payoff Convergence

DIFT's payoff for 100 perturbed strategies



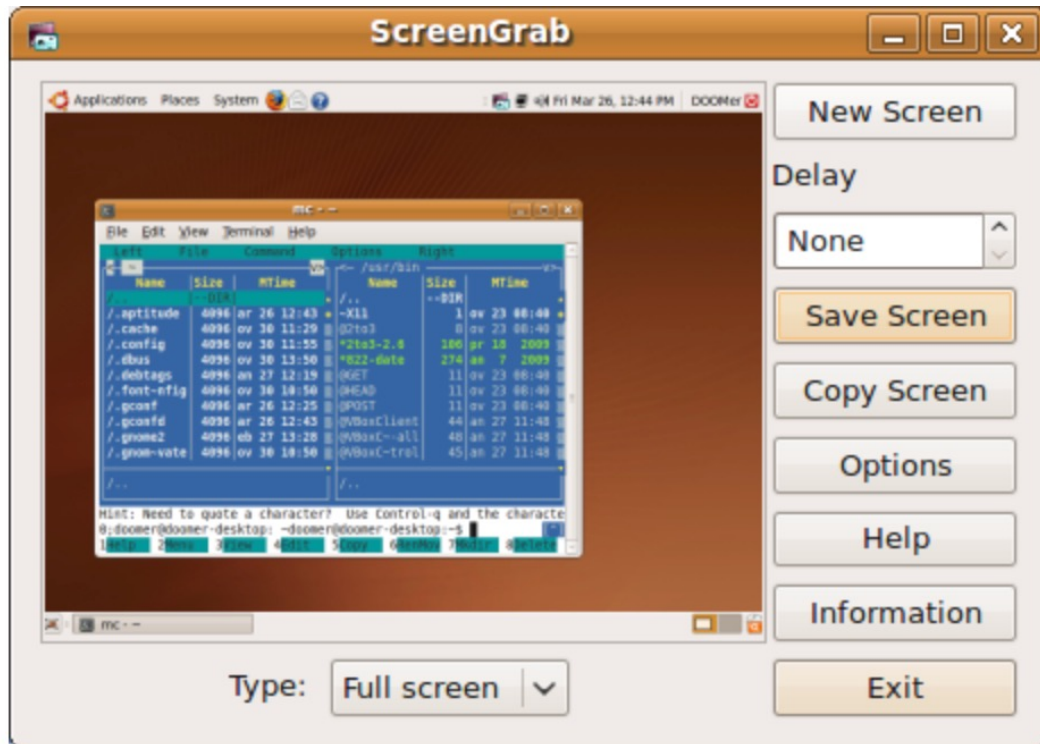
Convex Approximation

DIFT's payoff for 100 perturbed strategies



Original Function

Case Study: ScreenGrab Attack




- ScreenGrab: Linux program to take screenshots.
- Adversary wants to gain access to ScreenGrab process
- Fraction of flows traversing through each process ($Prob(s)$) extracted from RAIN log data



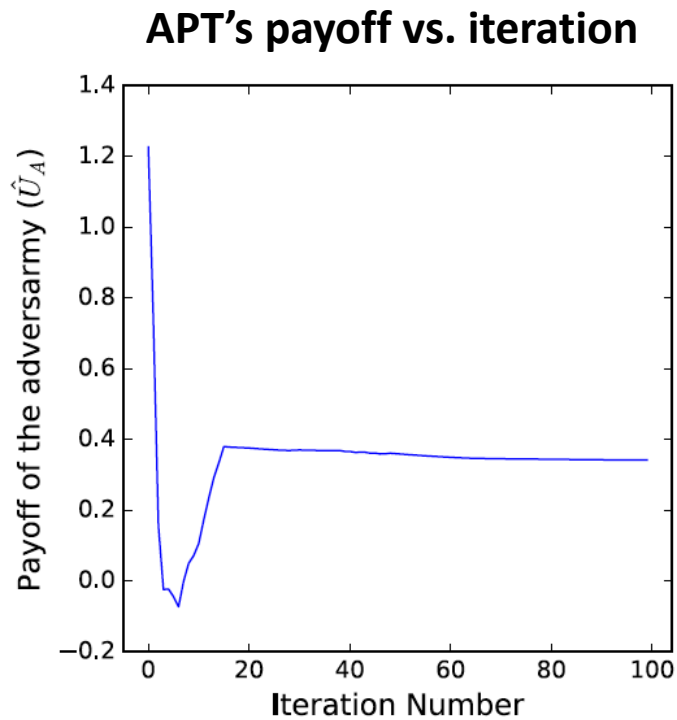
Case Study: ScreenGrab Attack

Pruned IFG for the ScreenGrab Attack

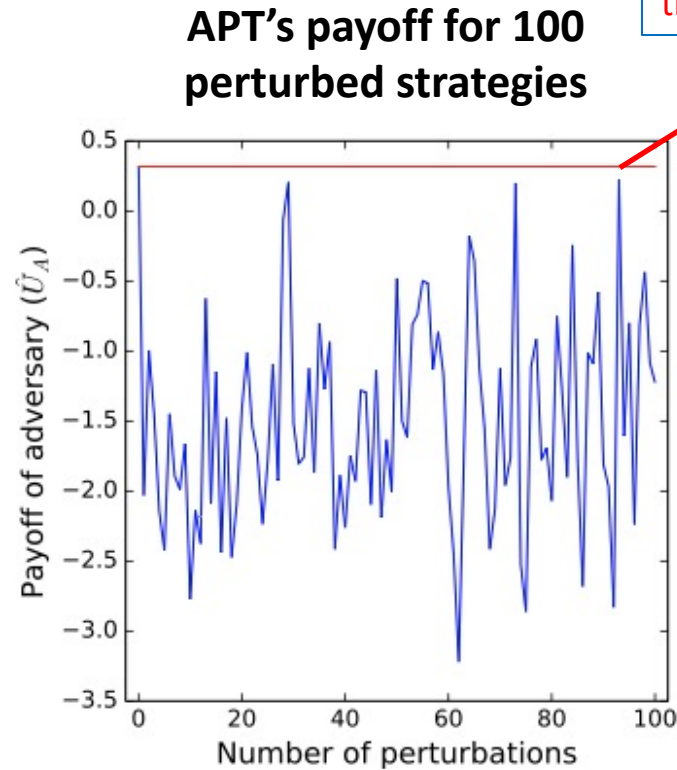


- ScreenGrab: Linux program to take screenshots.
- Adversary wants to gain access to ScreenGrab process (node 12)
- Fraction of flows traversing through each process ($Prob(s)$) extracted from RAIN log data

Simulation Results

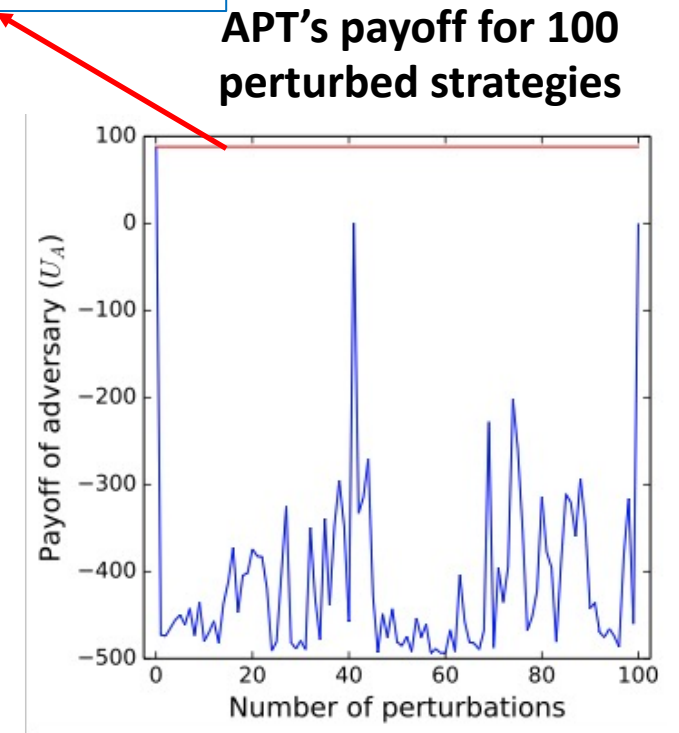


Payoff Convergence



Convex Approximation


Payoff received from at the ϵ -NE obtained from the algorithm



Original Function


Simulation Results

DIFT's payoff vs. iteration




Payoff Convergence

DIFT's payoff for 100 perturbed strategies



Convex Approximation

DIFT's payoff for 100 perturbed strategies



Original Function

Payoff received from at the ϵ -NE obtained from the algorithm



Observations

As can be seen, for most of the random perturbations, the payoff obtained by the non-neural network approach is indeed less than the payoff of the strategy returned by the neural network.

Specifically, in random graph experiment and for p_A , only two of the 100 perturbations resulted in higher payoff and for p_D , no perturbation yielded better payoff. Moreover, in ScreenGrab experiment, no perturbation resulted in higher payoff for either p_A or p_D . These results empirically validate that the proposed approach learns an approximate equilibrium.

Conclusion

- Formulated the interaction between APTs and DIFT as a two-player, multi-stage stochastic game with incomplete and imperfect information.
- Presented a supervised learning-based algorithm to learn an approximate Nash equilibrium for the game when the transition probabilities are unknown.

Future Work

- Characterize the convex approximation factor for the payoff functions of both players.
- Analyze the trade-off between obtaining a good convex approximation vs. the accuracy of the partial input convex neural networks.
- Investigate and apply the supervised learning approach to other types of games.

Thank You !

Email : shrm145@uw.edu

Project ADAPT website: <https://adapt.ece.uw.edu/>

